

CSE 678: Project

(<http://www.cse.ohio-state.edu/~kannan/678-wi12/proj/index.html>)

1 Overview

The goal of this project is to implement a TCP-like reliable transport layer protocol using the unreliable service provided by UDP, and then write a simple file transfer application to demonstrate its operation.

1. *Function Calls to be implemented*

The arguments and return values of these function calls must exactly match the ones for the corresponding function calls of UNIX socket implementation. To implement these functions, you can use any UDP related function calls.

- (a) *SOCKET*
- (b) *BIND*
- (c) *ACCEPT*
- (d) *CONNECT*
- (e) *SEND*
- (f) *RECV*
- (g) *CLOSE*

The focus of the project is on data transfer. Most of the TCP functionality will be implemented in the TCPD (TCP daemon) process which is equivalent to the TCP in the OS that runs in the background. These function calls will require communicating with the local TCPD process. The communication between the application process and the local TCPD process can be implemented with any inter-process communication mechanism such as UDP sockets. UDP communication within a machine can be assumed to be reliable.

Write a simple file-transfer application that uses your TCP implementation. Note that you will need this program for testing your TCP implementation. The file-transfer protocol will include a server called *ftps* and a client called *ftpc*. Start the server using the command

ftps <local-port>

Start *ftpc* with the command

ftpc <remote-IP> <remote-port> <local-file-to-transfer>

The *ftpc* client will send all the bytes of that local file using your implementation of TCP. The *ftps* server should receive the file and then store it. Make sure that after receiving the file at the *ftps* server you either give the file a different name or store it in a different directory than the original since all the CSE machines have your root directory mounted. Otherwise you will end up overwriting the original file.

The file-transfer application will use a simple format. The first 4 bytes (in network byte order) will contain the number of bytes in the file to follow. The next 20 bytes will contain the name of the file. The rest of the bytes to follow will contain the data in the file.

To simulate real network behavior, all communication between the two machines will go through local *troll* processes. *Troll* is a utility that allows you to introduce network losses and delay. More details on *troll* is provided later.

Here is some detail on the steps for transferring a file from machine M2 (client machine) to machine M1 (server machine).

- (a) Start the *troll* process and the TCPD process on machines M1 and M2.
- (b) On machine M1, start the file-transfer server *ftps*. It will make the function calls `SOCKET()`, `BIND()` and `ACCEPT()`. *ftps* will then block for a client to connect.
- (c) On machine M2, start the file-transfer client, *ftpc*. It will make function calls `SOCKET()`, `BIND()`, and `CONNECT()`.
- (d) Normally the `CONNECT()` should initiate TCP handshaking between the two TCPD processes. But in this project you are not implementing TCP handshaking. So `CONNECT()` is a null function.
- (e) The buffer management for this connection will be done in *TCPD_{M2}*.
- (f) *ftpc* will read bytes from the file and use the function `SEND()` to send data to *ftps*. The `SEND()` function call will need to send these bytes to the local TCPD process. The TCPD process will then store these bytes in a wrap-around buffer.
`SEND()` should be implemented as a blocking function call. It should not return until all bytes in the buffer passed in the argument is written in the TCPD buffer.
- (g) The buffer management functions will then take bytes from the buffer and create packets.
- (h) Upon receiving the first byte the TCPD on M1 will unblock the `ACCEPT()` call. The *ftps* application will then make calls to `RECV()` to receive data .
`RECV()` should not return until at least one byte is read from the TCPD buffer. However, if multiple contiguous bytes are available they should be read to fill up the buffer up to the maximum size specified in the argument.
- (i) After sending all the bytes of the file, *ftpc* closes the connection. The `CLOSE()` function call will initiate closing the connection.
- (j) On receiving all the bytes of the file, *ftps* will close the connection using the `CLOSE()` function.

2. *Connection Setup* As you are not implementing the TCP handshaking, the connection setup steps are fairly simple:

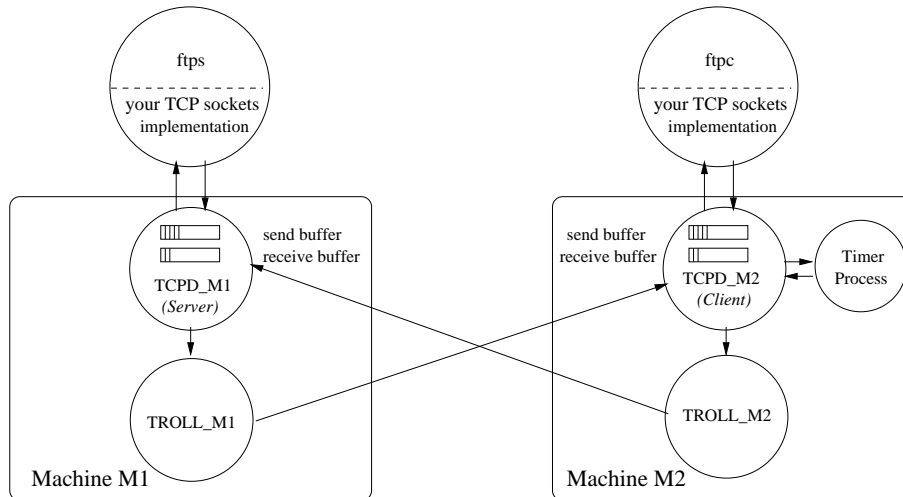


Figure 1: Connection Setup

- Initially the application process on M1 calls `ACCEPT()` and blocks. The `ACCEPT()` function in turn will send a message to the local TCPD and wait to hear back when a client has successfully connected.
- The `CONNECT()` function on M2 is an empty (or null) function.
- On reception of the first data packet, $TCPD_{M1}$ sends a message to the process waiting on `ACCEPT()`.

3. *RTT Computation*

Implement the Jacobson's algorithm (Section 21.3 of textbook) for computing RTT and RTO.

4. *Checksum Computation*

The CRC (Cyclic Redundancy Code) checksumming technique should be used for computing the checksum.

5. *Packet Formats*

The TCP packet structure should be strictly followed for both TCP and ACK packets. Instead of TCP's cumulative ACK, the ACK packet will acknowledge the data packet just received. Note that each packet will be ACKed.

6. *Timer Implementation*

Each data packet after transmission will require a timer to be started. When the timer runs out, the packet will need to be re-transmitted. Since a large number of packets may be in transit at any given time, a large number of timers may be simultaneously running.

Instead of using explicit timers for each packet, you will implement the timers using a delta-list. More details on delta-list is available on the project web-site.

The delta list must be maintained in a separate process called the "timer-process". When a new timer needs to be started, a message is sent to the local timer process, indicating how long the timer needs to run for, which port the timer process should send notification upon expiry, and the byte sequence number of the packet for which this timer is being started.

7. *Buffer Management and Sliding Window Protocol*

Implement the selective repeat algorithm. Use a fixed window size of 20. You are not required to implement slow-start, congestion control, or flow control algorithms.

The send and receive buffers will be wrap-around (or circular) buffers. Use a buffer of 64 KB for sending and 64 KB for receiving and MSS of 1000 Bytes.

Buffers/arrays in other processes should not store more than 1 MSS worth of data.

8. *Connection Shutdown*

Implement the exact state diagram for shutdown of a TCP connection. All the data structures related to the socket will be deallocated. However the buffer management function should make sure that all data has been acknowledged before deallocating the data structures.

2 Troll

In the CSE network it is hard to artificially create real network scenarios (lossy links, packet garbling etc.) Use the *troll* utility to control the rate of garbling, discarding, delaying or duplication of packets. All packets will first go through a local troll process running on the same machine, where they will be subject to delay, garbling and/or drops. The packets will then be forwarded to the intended destination. The final destination should be marked in the first 16-bytes of the packet sent to the troll process. (Read the *troll* manual)




Figure 2: Troll

Test the *troll* program using the *totroll* and *fromtroll* programs. Source code for the two programs is available on the project web page. Here is an example of how you can test the troll functionality:

- On machine A, start *troll* to communicate on port 10001 using the command:
troll 10001

- On machine B, start *fromtroll* to listen on port 10002 using the command

fromtroll 10002

- On machine A, use *totroll* to send a short message via *troll* to *fromtroll* on machine B using the following command.

totroll A 10001 B 10002

A and B have to be replaced by the IP addresses of the corresponding machines. You can use a combination of the following commands to find the IP address of your machine: *nslookup*, *hostname*, *ifconfig*. You can choose any two machines to run the *fromtroll* and *totroll* programs. A list of CSE machines can be obtained by using a command such as *rusers*. Note that the particular port numbers used in the example above may be unavailable if some other process is using it. To login to a different machine, you may use *ssh*, *telnet* or *rlogin*.

3 Testing and Developing the Project

This project has several components. Debug and test each function carefully before integrating it into your code. Start your work very early as certain components such as the sliding window protocol may take a long time to debug.

4 Checkpoints

- **Due Jan 5 (Tuesday), in class, Group Formations:** Write your and your partner's name in the signup sheet to be passed in class.
- **Due Jan 19 (Thursday), in class, Project Proposal:** The proposal must include details on the wrap around buffer, packet formats (formats for all packets exchanged between TCPD, FTPC, FTPS, TROLL, and TIMER processes), timer process, checksumming algorithm, RTT/RTO computation, description of operations within the implementation of all CAPITAL functions, and connection shutdown. Details of various data structures and when/how

they are updated must be included for each of those modules. In addition, present a clear timeline and work distribution plan with specific internal milestones. Proper planning with time allocation for debugging and testing is crucial for successful execution of the project.

It must be typeset using a formatting tool, such as Word, Framemaker, Latex etc. Use 11pt font with single spacing and single column. The recommended length is 6-8 pages.

- **Feb 14th/16th (Tuesday/Thursday) Checkpoint Demo and Submission:** Be creative in how you demonstrate the correct functionality of the modules (see table).

In addition to the demo, you need to submit your code using the *submit* utility by 11:59pm, Feb 16th. Use the following command for submitting the code :

submit c678aa lab5 <code-directory-name>

Your code directory must contain a README file that describes all the C files in your directory. Also indicate how to test the functionality of each module. It must contain a Makefile. If you resubmit the files, submit both the code and the final report again, as each invocation of *submit* deletes all files of the previous submission for the same lab. Read *man submit* for clarification. Only one person from each group must submit the project.

- **Mar 6 (Tuesday), Final Demo:** The final demonstrations will be held in Room: DL280 Time: 4:30-7:30pm. There will be a sign-up sheet for each team.
- **Mar 8 (Thursday), 11:59pm EST, Final Report and Code:** Submit the final version of your code and the final project report by this deadline. The final project report should include details on the implementation, discussion on all the features of your program (including any extra features beyond what is required), optimizations, and possible enhancements for the future. The report could include code fragments only if it is absolutely essential.

The report must be typeset using a formatting tool, such as Word, Framemaker, Latex etc. Use 11pt font with single spacing. The report must be 6-10 pages long.

Submit your code and final report using the *submit* utility. Use the following command for submitting the final project code and the final report:

submit c678aa lab6 <code-directory-name> <final-report-file-name>

Your code directory must contain a README file that describes all the C files in your directory. Also indicate how to run your program. It must contain a Makefile. If you resubmit the files, submit both the code and the final report again, as each invocation of *submit* deletes all files of the previous submission for the same lab. Read *man submit* for clarification. Only one person from each group must submit the project.

5 Miscellaneous

- **Team Formation:** You can either work by yourself or in a group of 2 (recommended). To the extent possible, please make sure that your group partner is a motivated and hard-working student. A significant portion of the grade depends on the project, and both the team members will get equal scores in the project.
- **Platform:** Use the *stdsun* system (Not the testbed) for all your implementation.
- **Delayed Submissions:** Late demonstrations, code submissions or report submissions are not eligible for any points.

Deadline	Item	details	Points
In class Jan 19	Proposal		6
	Wrap around buffer	1	
	Packet Formats (formats for all packets exchanged between TCPD, FTPC, FTPS, TROLL, and TIMER processes)	1	
	Timer process	1	
	Checksumming Algorithm	1	
	RTT and RTO	0.5	
	Description of operations within the implementation of all CAPITAL functions	0.5	
	Connection Shutdown	0.5	
	Timeline and Work Distribution	0.5	
DL266 Feb 14/16 (3:30-5:30)	Checkpoint Demonstration		14
	Delta-timer	7	
	Checksumming	7	
Submit 11:59pm, Feb 16	Well-documented code		2
DL280 Mar 6 (4:30-7:30pm)	Final Demonstration		22
	RTT computation	4	
	Circular Buffer management	6	
	Shutdown	2	
	Successful execution with troll on both machines (with garble 25%, destroy 25%, duplicate 25%, reorder, exponential delay with mean 10ms)	10	
Submit 11:59pm Mar 8	Final Report		4
	Project overview and details on each process	1	
	Detailed description of each project component (Wraparound buffer, TIMER, checksumming, RTT/RTO, Packet formats, Implementation of SOCKET, BIND, SEND, RECV, CLOSE)	1	
	Detailed description of how to compile, run and test the program	1	
	Possible future extensions to make this program more efficient and to add more features	1	
Submit 11:59pm Mar 8	Well-documented code		2
	TOTAL		50